

# Mini Project 1: Parallel LU Decomposition

20200703 SOONHO KIM

## 1. Program Structure

### Matrix Allocation:

The matrix is represented as a 2D dynamic array (double\*\*) and allocated via a helper function `allocate_matrix(n)`, which returns an n x n matrix of 64-bit floating-point values.

### Random Initialization:

Matrix elements are initialized with uniformly distributed random numbers in [0, 1], using C++11's `<random>` library. To support parallel generation, each thread uses its own instance of `mt19937`, seeded with `std::mt19937 rng(time(NULL) + tid);`

This ensures thread-safe and reproducible randomness across threads under OpenMP.

### LU Decomposition Function:

The decomposition follows handout pseudocode and includes the following steps:

1. Pivot Selection:  
Find the row with the maximum absolute value in the current column to enhance numerical stability.
2. Row Swapping:  
Swap the current row with the pivot row in both the matrix and the permutation vector.
3. Multiplier Computation:  
Compute values for the lower triangular matrix LLL using:
$$L(i, k) = \frac{A(i, k)}{U(k, k)}$$
4. Submatrix Update (Schur Complement):  
Update the remaining submatrix entries with

$$A(i, j) = A(i, j) - L(i, k) * U(k, j)$$

This step reduces the original matrix to upper and lower triangular form in-place.

## 2. Parallelization Strategy

The LU decomposition is parallelized using OpenMP, primarily through for-based directives to leverage loop-level parallelism.

- **Parallel Directives Used:**
  1. `#pragma omp parallel for`: for standard loop parallelization.
  2. `#pragma omp parallel for collapse(2)`: for nested loop optimization, particularly when updating submatrices.
- **Parallelization Points:**
  1. Matrix Initialization in `main()`: threads generate random values in parallel.
  2. L Calculation Loop in `lu_decomposition()`: each row update is independent.
  3. Schur Complement Update in `lu_decomposition()`: nested loops updating submatrix elements are parallelized.
- **Data Partitioning and Synchronization:**

Since all parallel regions are loop-based, the workload is implicitly partitioned along i and j loop indices. There are no shared resource conflicts due to row-wise memory access, and false sharing is minimized.

### 3. Experimental Setup

The program was tested under two hardware environments using the same code:

- Test Platforms:
  - AMD Ryzen 5 5600G (6 cores / 12 threads, 16MB L3 cache)
  - Intel Core i7-1360P (4 performance + 8 efficient cores, 16 threads, 18MB L3 cache)
- Parameters:
  - Matrix sizes:  $N \in \{1000, 2000, 4000, 8000\}$
  - Number of threads:  $p \in \{1, 2, 4, 8, 12, 16\}$
  - Each configuration was repeated 5 times and averaged.
- Metrics Collected:
  - Execution time  $T(p)$
  - Residual  $\|PA - LU\|_{2,1}$

### 4. Performance Evaluation

Before measuring all performance data, I first verified the correctness of the LU decomposition as instructed in the handout.

Specifically, I computed the **L2,1 norm** of the residual matrix, defined as the sum of Euclidean norms of the columns of  $PA - LU$ . The result showed **very small values in the range of  $10^{-10} \sim 10^{-11}$** , which confirmed the numerical correctness.

Based on this, I proceeded to measure the LU decomposition time across various test cases.

| MatrixSize | Threads | Type     | Time(s)   | Residual | MatrixSize | Threads | Type     | Time(s)  | Residual |
|------------|---------|----------|-----------|----------|------------|---------|----------|----------|----------|
| 1000       | 1       | Serial   | 0.152936  | 4.17E-11 | 1000       | 8       | Parallel | 0.053539 | 4.19E-11 |
| 1000       | 1       | Serial   | 0.142234  | 4.24E-11 | 1000       | 8       | Parallel | 0.054411 | 4.19E-11 |
| 1000       | 1       | Serial   | 0.146328  | 4.21E-11 | 1000       | 8       | Parallel | 0.055908 | 4.16E-11 |
| 1000       | 1       | Serial   | 0.143175  | 4.18E-11 | 1000       | 8       | Parallel | 0.063236 | 4.16E-11 |
| 1000       | 1       | Parallel | 0.181118  | 4.18E-11 | 1000       | 12      | Parallel | 0.353676 | 4.20E-11 |
| 1000       | 1       | Parallel | 0.18023   | 4.13E-11 | 1000       | 12      | Parallel | 0.124319 | 4.16E-11 |
| 1000       | 1       | Parallel | 0.189719  | 4.16E-11 | 1000       | 12      | Parallel | 0.109523 | 4.19E-11 |
| 1000       | 1       | Parallel | 0.188166  | 4.19E-11 | 1000       | 12      | Parallel | 0.049384 | 4.19E-11 |
| 1000       | 2       | Parallel | 0.100916  | 4.15E-11 | 2000       | 8       | Parallel | 1.26949  | 2.03E-10 |
| 1000       | 2       | Parallel | 0.0952335 | 4.14E-11 | 2000       | 8       | Parallel | 1.69776  | 2.03E-10 |
| 1000       | 2       | Parallel | 0.0960171 | 4.14E-11 | 2000       | 8       | Parallel | 1.50775  | 2.04E-10 |
| 1000       | 2       | Parallel | 0.0926738 | 4.15E-11 | 2000       | 8       | Parallel | 1.66659  | 2.04E-10 |
| 1000       | 4       | Parallel | 0.0818114 | 4.20E-11 | 2000       | 12      | Parallel | 1.58414  | 2.03E-10 |
| 1000       | 4       | Parallel | 0.0692526 | 4.20E-11 | 2000       | 12      | Parallel | 2.10347  | 2.03E-10 |
| 1000       | 4       | Parallel | 0.0716119 | 4.20E-11 | 2000       | 12      | Parallel | 1.93013  | 2.03E-10 |
| 1000       | 4       | Parallel | 0.0797535 | 4.23E-11 | 2000       | 12      | Parallel | 1.34326  | 2.04E-10 |
| 2000       | 1       | Serial   | 2.5869    | 2.04E-10 |            |         |          |          |          |
| 2000       | 1       | Serial   | 2.55105   | 2.04E-10 |            |         |          |          |          |
| 2000       | 1       | Serial   | 2.22259   | 2.05E-10 |            |         |          |          |          |
| 2000       | 1       | Serial   | 2.46416   | 2.05E-10 |            |         |          |          |          |
| 2000       | 1       | Parallel | 2.9385    | 2.03E-10 |            |         |          |          |          |
| 2000       | 1       | Parallel | 2.71701   | 2.03E-10 |            |         |          |          |          |
| 2000       | 1       | Parallel | 2.59234   | 2.03E-10 |            |         |          |          |          |
| 2000       | 1       | Parallel | 2.51985   | 2.02E-10 |            |         |          |          |          |
| 2000       | 2       | Parallel | 1.9152    | 2.04E-10 |            |         |          |          |          |
| 2000       | 2       | Parallel | 2.17923   | 2.03E-10 |            |         |          |          |          |
| 2000       | 2       | Parallel | 1.88516   | 2.04E-10 |            |         |          |          |          |
| 2000       | 2       | Parallel | 1.7937    | 2.04E-10 |            |         |          |          |          |
| 2000       | 4       | Parallel | 1.48389   | 2.03E-10 |            |         |          |          |          |
| 2000       | 4       | Parallel | 1.73146   | 2.03E-10 |            |         |          |          |          |
| 2000       | 4       | Parallel | 1.67667   | 2.04E-10 |            |         |          |          |          |
| 2000       | 4       | Parallel | 1.63847   | 2.04E-10 |            |         |          |          |          |

Table1. LU decomposition Efficiency


| p  |          | 1        |       |      |          | 2     |      |          |       | 4    |          |       |      | 8        |       |      |         | 12    |      |      |       | 16  |  |  |  |
|--|----------|----------|-------|------|----------|-------|------|----------|-------|------|----------|-------|------|----------|-------|------|---------|-------|------|------|-------|-----|--|--|--|
| N  | S        | T(p)     | spdUp | pEf  | T(p)     | spdUp | pEf  | T(p)     | spdUp | pEf  | T(p)     | spdUp | pEf  | T(p)     | spdUp | pEf  | T(p)    | spdUp | pEf  | T(p) | spdUp | pEf |  |  |  |
| AMD Ryzen 5 5600G (specifications: 6P/12Threads, L3 cache 16MB)                      |          |          |       |      |          |       |      |          |       |      |          |       |      |          |       |      |         |       |      |      |       |     |  |  |  |
| 1000   | 0.1462   | 0.1848   | 0.8   | 0.79 | 0.0962   | 1.5   | 0.76 | 0.0756   | 1.9   | 0.48 | 0.0568   | 2.6   | 0.32 | 0.1592   | 0.9   | 0.08 | -       | -     | -    | -    | -     |     |  |  |  |
| 2000   | 2.4562   | 2.6919   | 0.9   | 0.91 | 1.9433   | 1.3   | 0.63 | 1.6326   | 1.5   | 0.38 | 1.5354   | 1.6   | 0.20 | 1.7400   | 1.4   | 0.12 | -       | -     | -    | -    | -     |     |  |  |  |
| 4000   | 23.3884  | 24.4122  | 1.0   | 0.96 | 19.6104  | 1.2   | 0.60 | 19.3466  | 1.2   | 0.30 | 22.1187  | 1.1   | 0.13 | 21.0237  | 1.1   | 0.09 | -       | -     | -    | -    | -     |     |  |  |  |
| 8000   | 179.0600 | 190.4000 | 0.9   | 0.94 | 155.3750 | 1.2   | 0.58 | 158.0830 | 1.1   | 0.28 | 169.5430 | 1.1   | 0.13 | 180.6020 | 1.0   | 0.08 | -       | -     | -    | -    | -     |     |  |  |  |
| 13th Gen Intel(R) Core(TM) i7-1360P (specifications: 4P+8E/16Threads, L3 cache 18MB) |          |          |       |      |          |       |      |          |       |      |          |       |      |          |       |      |         |       |      |      |       |     |  |  |  |
| 1000   | 0.2747   | 0.4764   | 0.6   | 0.58 | 0.2603   | 1.1   | 0.53 | 0.1909   | 1.4   | 0.36 | 0.2640   | 1.0   | 0.13 | 0.8486   | 0.3   | 0.03 | 1.2225  | 0.2   | 0.01 | -    | -     |     |  |  |  |
| 2000   | 3.2924   | 3.8763   | 0.8   | 0.85 | 2.1246   | 1.5   | 0.77 | 1.3644   | 2.4   | 0.60 | 0.7708   | 4.3   | 0.53 | 1.3588   | 2.4   | 0.20 | 1.4993  | 2.2   | 0.14 | -    | -     |     |  |  |  |
| 4000   | 21.5544  | 28.8771  | 0.7   | 0.75 | 16.5034  | 1.3   | 0.65 | 11.0607  | 1.9   | 0.49 | 6.7043   | 3.2   | 0.40 | 8.7390   | 2.5   | 0.21 | 10.1280 | 2.1   | 0.13 | -    | -     |     |  |  |  |
| 8000   | 174.8510 | 224.3920 | 0.8   | 0.78 | 127.6360 | 1.4   | 0.68 | 82.8351  | 2.1   | 0.53 | 51.1176  | 3.4   | 0.43 | 67.2453  | 2.6   | 0.22 | 74.0107 | 2.4   | 0.15 | -    | -     |     |  |  |  |

The data I measured are as follows. In consideration of the number of cores and threads of each CPU, Ryzen measured only up to p=12.

All graphs and data described later are based on the data in Table 1.

**Table2. LU Decomposition Time Comparison** (for each matrix size N and thread count p)

|          | S       | T(1)    | T(2)    | T(4)    | T(8)    | T(12)   |
|----------|---------|---------|---------|---------|---------|---------|
| 1000 (R) | 0.1462  | 0.1848  | 0.0962  | 0.0756  | 0.0568  | 0.1592  |
| 2000 (R) | 2.4562  | 2.6919  | 1.9433  | 1.6326  | 1.5354  | 1.7400  |
| 4000 (R) | 23.3884 | 24.4122 | 19.6104 | 19.3466 | 22.1187 | 21.0237 |
| 1000 (I) | 0.2747  | 0.4764  | 0.2603  | 0.1909  | 0.2640  | 0.8486  |
| 2000 (I) | 3.2924  | 3.8763  | 2.1246  | 1.3644  | 0.7708  | 1.3588  |
| 4000 (I) | 21.5544 | 28.8771 | 16.5034 | 11.0607 | 6.7043  | 8.7390  |




The left column indicates the hardware used: **R (Ryzen)** and **I (Intel)**, along with the corresponding matrix size N. The horizontal entries show **S (Serial time)** and **T(p)** in seconds.

On Ryzen, the time reduction from parallel execution is **less significant**, showing limited performance gain through multi-threading.

**Table3. SpeedUp vs. Number of Threads**

|          | T(1) | T(2) | T(4) | T(8) | T(12) | T(16) |
|----------|------|------|------|------|-------|-------|
| 1000 (R) | 0.8  | 1.5  | 1.9  | 2.6  | 0.9   |       |
| 2000 (R) | 0.9  | 1.3  | 1.5  | 1.6  | 1.4   |       |
| 4000 (R) | 1.0  | 1.2  | 1.2  | 1.1  | 1.1   |       |
| 8000 (R) | 0.9  | 1.2  | 1.1  | 1.1  | 1.0   |       |
| 1000 (I) | 0.6  | 1.1  | 1.4  | 1.0  | 0.3   | 0.2   |
| 2000 (I) | 0.8  | 1.5  | 2.4  | 4.3  | 2.4   | 2.2   |
| 4000 (I) | 0.7  | 1.3  | 1.9  | 3.2  | 2.5   | 2.1   |
| 8000 (I) | 0.8  | 1.4  | 2.1  | 3.4  | 2.6   | 2.4   |



From this point on, I analyze the results based on the Intel CPU.


The data shows that for matrix sizes up to **N = 1000**, the best speedup is achieved when the number of threads is **p = 2**.

When the matrix size exceeds **N = 2000**, the maximum speedup occurs at **p = 8**.

However, increasing the thread count beyond this point actually **degrades performance**, likely due to synchronization overhead and limited parallelism benefits.

**Table4. Parallel Efficiency vs. Number of Threads**

|          | T(1) | T(2) | T(4) | T(8) | T(12) | T(16) |
|----------|------|------|------|------|-------|-------|
| 1000 (R) | 0.79 | 0.76 | 0.48 | 0.32 | 0.08  |       |
| 2000 (R) | 0.91 | 0.63 | 0.38 | 0.20 | 0.12  |       |
| 4000 (R) | 0.96 | 0.60 | 0.30 | 0.13 | 0.09  |       |
| 8000 (R) | 0.94 | 0.58 | 0.28 | 0.13 | 0.08  |       |
| 1000 (I) | 0.58 | 0.53 | 0.36 | 0.13 | 0.03  | 0.01  |
| 2000 (I) | 0.85 | 0.77 | 0.60 | 0.53 | 0.20  | 0.14  |
| 4000 (I) | 0.75 | 0.65 | 0.49 | 0.40 | 0.21  | 0.13  |
| 8000 (I) | 0.78 | 0.68 | 0.53 | 0.43 | 0.22  | 0.15  |



Parallel efficiency indicates how effectively each thread contributes to performance.

As the number of threads **p** increases, **E(p)** drops sharply, indicating diminishing returns due to overhead.

In particular, when **p** exceeds 8, the efficiency falls below 0.3, suggesting that the overhead of parallelization becomes significant and outweighs the performance gains.

## 5. Analysis & Discussion

- The results show that larger matrix sizes generally lead to better parallel performance. In my environment, significant speedup was observed when **N > 2000**.
- Parallel efficiency decreases** as the number of threads **p** increases. This is likely due to increased communication overhead and cache contention.
- The hardware architecture has a substantial impact on parallel performance:
  - AMD Ryzen** uses regular cores, while **Intel i7-1360P** employs a hybrid architecture (Performance + Efficient cores), which seems to provide better scalability in parallel workloads.
- The experiments confirmed a typical parallel trade-off:
  - Speedup increases** with more threads, but **efficiency tends to decrease** due to overhead, especially beyond a certain thread count.

## 6. Conclusion

- LU decomposition with OpenMP shows **limited efficiency for small matrix sizes**.
- The parallel implementation is **effective up to 8 threads**.
- Future optimization directions** include:
  - Block-based decomposition to improve data locality.
  - Cache-aware memory access strategies.
  - Exploiting NUMA-aware memory allocation** (e.g., using `libnuma` or `numactl`) to reduce memory latency and improve performance in multi-socket environments.